

FUSE

(Filesystem in Userspace)

Krerik Piromsopa, Ph.D.

Why Userspace?

- Writing (good) code is not easy.

Let's speed it up....

- Writing (good) code in kernel is hard

- No libc (no printf, Qt, stdio...)

- Too many reboots, Kernel Panic

Building a filesystem is difficult?

- Performance reduce disk seek
- Life cycle of disk (SSD block limited life cycle)
- Maximum filesize limitations
- Metadata, permission, security (encryption)
- Deduplication, compression, snapshots
- etc....

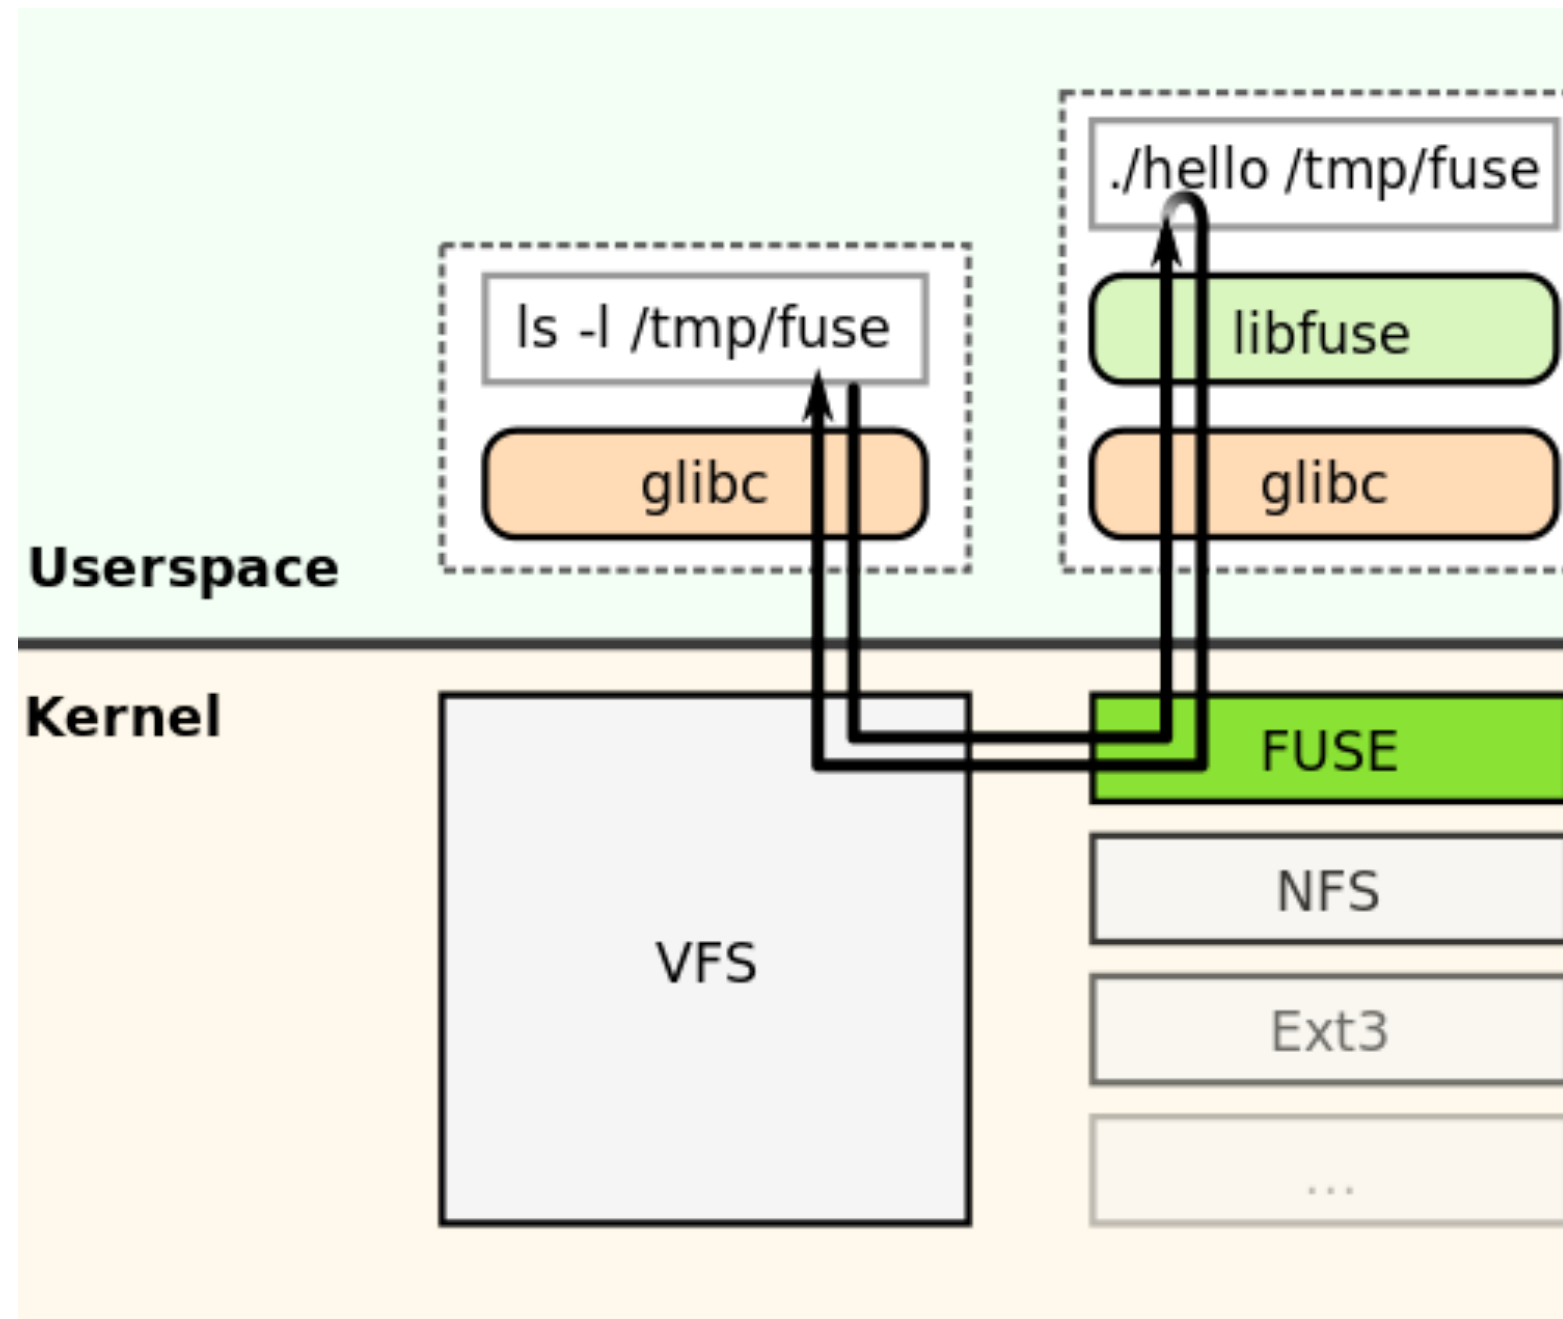
What is FUSE?

- Writing a filesystem in
- your favorite environments
(language, library, debugging tool)

running file system code in user space while the FUSE module provides only a "bridge" to the actual kernel interfaces. [wikipedia]

How does it work

- File/Filesystem API are passed to a process.
- A kernel module (driver) is a bridge
- Available on Linux, BSD (MacOSX), Android, Minix, ...
- Doken is a Windows implementation (incomplete?)



binding for C, Python, Objective-C, Ruby, Java, C#,.....

Operations (1)

- `int (*getattr) (const char *, struct stat *);`
 - Get file attributes.
- `int (*readlink) (const char *, char *, size_t);`
 - Read the target of a symbolic link
- `int (*mknod) (const char *, mode_t, dev_t);`
 - Create a file node.
- `int (*mkdir) (const char *, mode_t);`
 - Create a directory. Note that the mode argument may not have the type specification bits set, i.e. `S_ISDIR(mode)` can be false. To obtain the correct directory type bits use `mode | S_IFDIR`

https://lastlog.de/misc/fuse-doc/doc/html/structfuse__operations.html

Operations (2)

- `int (*unlink) (const char *)`;
 - Remove a file
- `int (*rmdir) (const char *)`;
 - Remove a directory
- `int (*symlink) (const char *, const char *)`;
 - Create a symbolic link
- `int (*rename) (const char *, const char *)`;
 - Rename a file
- `int (*link) (const char *, const char *)`;
 - Create a hard link to a file
- `int (*chmod) (const char *, mode_t)`;
 - Change the permission bits of a file
- `int (*chown) (const char *, uid_t, gid_t)`;
 - Change the owner and group of a file
- `int (*truncate) (const char *, off_t)`;
 - Change the size of a file
- `int (*open) (const char *, struct fuse_file_info *)`;
 - File open operation.
-

https://lastlog.de/misc/fuse-doc/doc/html/structfuse__operations.html

Operations (3)

- `int (*read) (const char *, char *, size_t, off_t, struct fuse_file_info *);`
 - Read data from an open file.
- `int (*write) (const char *, const char *, size_t, off_t, struct fuse_file_info *);`
 - Write data to an open file
- `int (*statfs) (const char *, struct statvfs *);`
 - Get file system statistics
- `int (*flush) (const char *, struct fuse_file_info *);`
 - Possibly flush cached data

Operations (4)

- `int (*opendir) (const char *, struct fuse_file_info *)`;
 - Open directory. Unless the 'default_permissions' mount option is given, this method should check if `opendir` is permitted for this directory. Optionally `opendir` may also return an arbitrary filehandle in the `fuse_file_info` structure, which will be passed to `readdir`, `closedir` and `fsyncdir`.
- `int (*readdir) (const char *, void *, fuse_fill_dir_t, off_t, struct fuse_file_info *)`;
 - Read directory
- `int (*releasedir) (const char *, struct fuse_file_info *)`;
 - Release directory
- `int (*fsyncdir) (const char *, int, struct fuse_file_info *)`;
 - Synchronize directory contents
- `void *(*init) (struct fuse_conn_info *conn)`;
 - Initialize file system.
- `int (*lock) (const char *, struct fuse_file_info *, int cmd, struct flock *)`;
 - Perform POSIX file locking operation

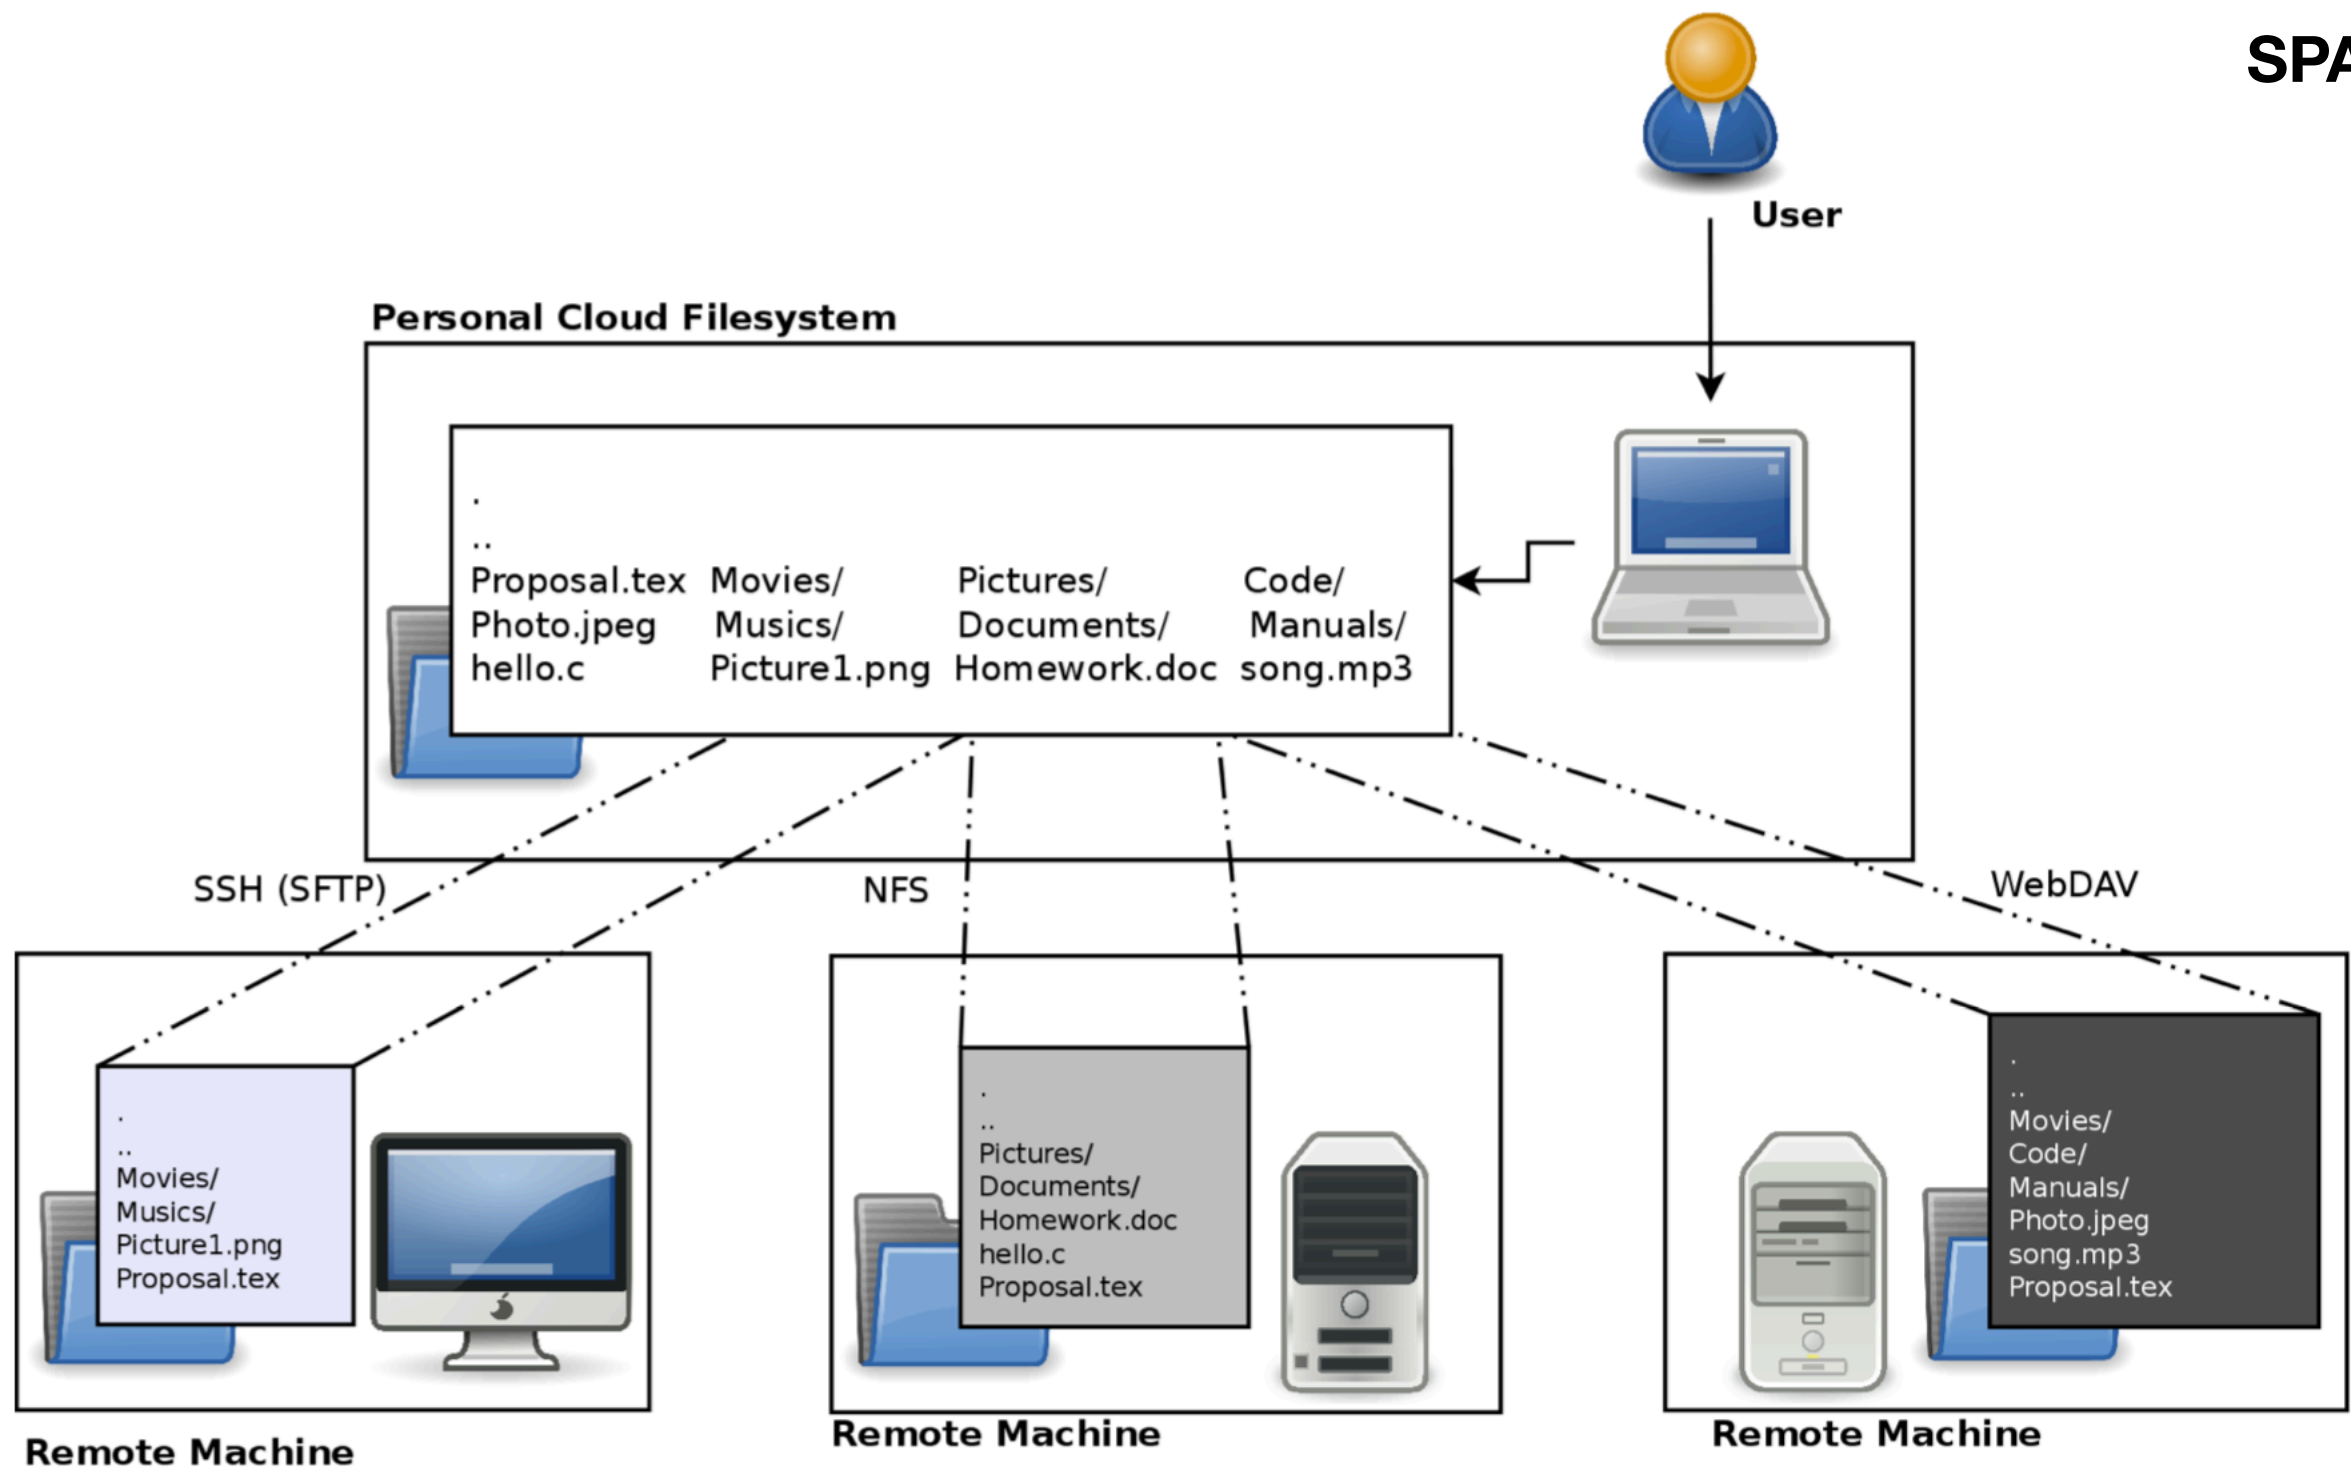
Some Implementations

- GmailFS
- EncFS
- NTFS-3G
- WikipediaFS
- Lustre
- SSHFS
- FTPFS
- ImapFS
- YoutubeFS
- Gdrive, Grive?
- gitFS
- etc... (build your own)

Some from my collections

- S. Dhumbumroong and K. Piromsopa, “Personal Cloud Filesystem: A distributed unification filesystem for personal computer and portable device,” in Computer Science and Software Engineering (JCSSE), 2011 Eighth International Joint Conference on, 2011, pp. 58–62.
(for both Linux and Windows) **SPAFS**
- W. Ratinimitum and K. Piromsopa, “An implementation of RESTful-based Scalable File System,” in JCSSE 2012 - 9th International Joint Conference on Computer Science and Software Engineering, 2012, pp. 136–141.
- More from students’ projects in the past

SPAFS



Python-FUSE API

- `create(path, mode)`
- `truncate(path, size)`
- `mknod(path, mode, dev)`
- `open(path, mode)`
- `write(path, data, offset)`
- `read(path, length, offset)`
- `release(path)`
- `fsync(path)`
- `chmod(path, mode)`
- `chown(path, oid, gid)`
- `mkdir(path, mode)`
- `unlink(path)`
- `readdir(path)`
- `rmdir(path)`
- `rename(opath, npath)`
- `ink(srcpath, dstpath)`

File Operations and API

Reading
(e.g. cat a.txt)

Writing
(e.g. echo "a" > a.txt)

Appending
(e.g. echo "a" >> a.txt)

Truncating
(e.g. echo "a" > a.txt)

getattr

open

read

release

getattr

create

write

flush

release

getattr

open

write

flush

release

getattr

truncate

open

write

flush

release

Removing
(e.g. rm a.txt)

getattr

unlink

Hint... Try strace

File Operations and API

Creating Directory
(e.g. mkdir demo)

getattr

mkdir

Reading Directory
(e.g. ls demo/)

getattr

opendir

readdir

releasedir

Removing Directory
(e.g. rmdir demo)

getattr

rmdir

Changing Permission
(e.g. chmod 777 demo)

getattr

chmod

Changing Ownership
(e.g. chmod 777 demo)

getattr

chown

Linking
(e.g. ln -s demo d)

getattr

symlink

Renaming
(e.g. mv demo d)

getattr

rename

Let's see a demo

**Let's rethink and build
your own filesystem**