

*.0

Krerik Piromsopa, Ph. D.

What is Kernel Module?

- Pieces of code that can be loaded and unloaded into the kernel upon demand.
- Extend the functionality of the kernel without rebooting the system.
 - Hardware Drivers
 - Special API
 - etc.

Linux kernel diagram

functions

system

networking

storage

memory

processing

human
interface

layers

user space
interfaces

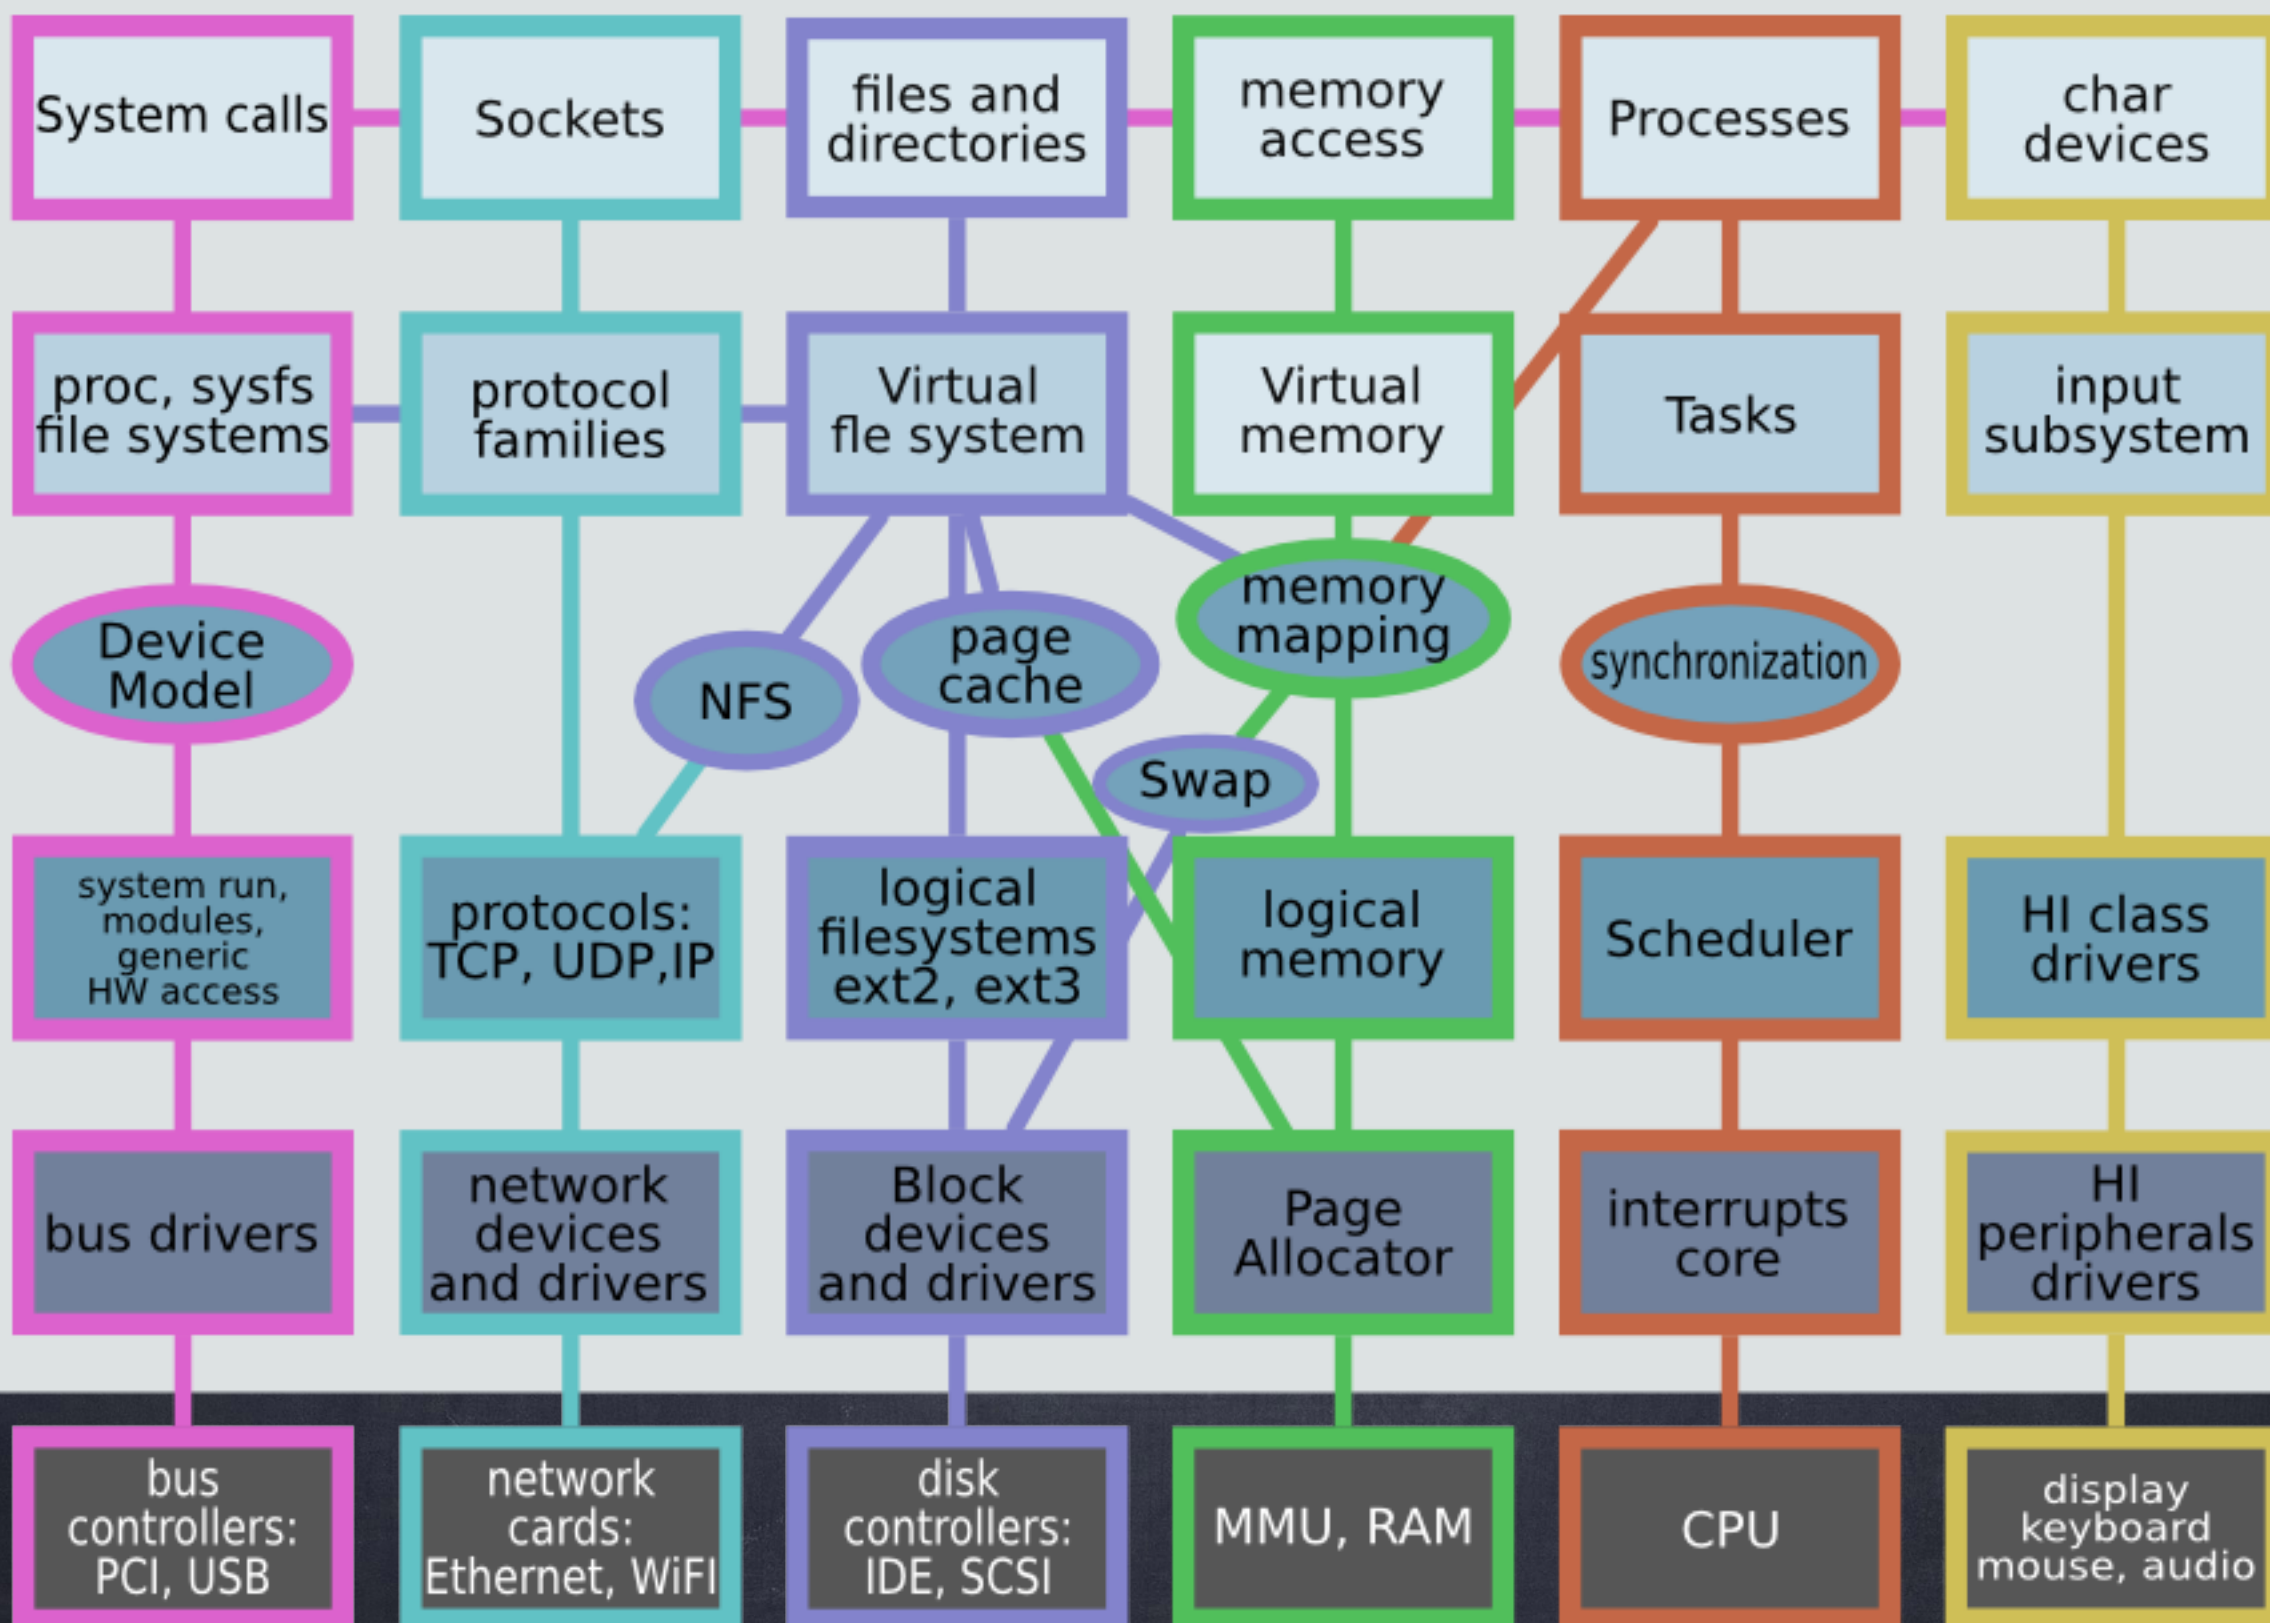
virtual
subsystems

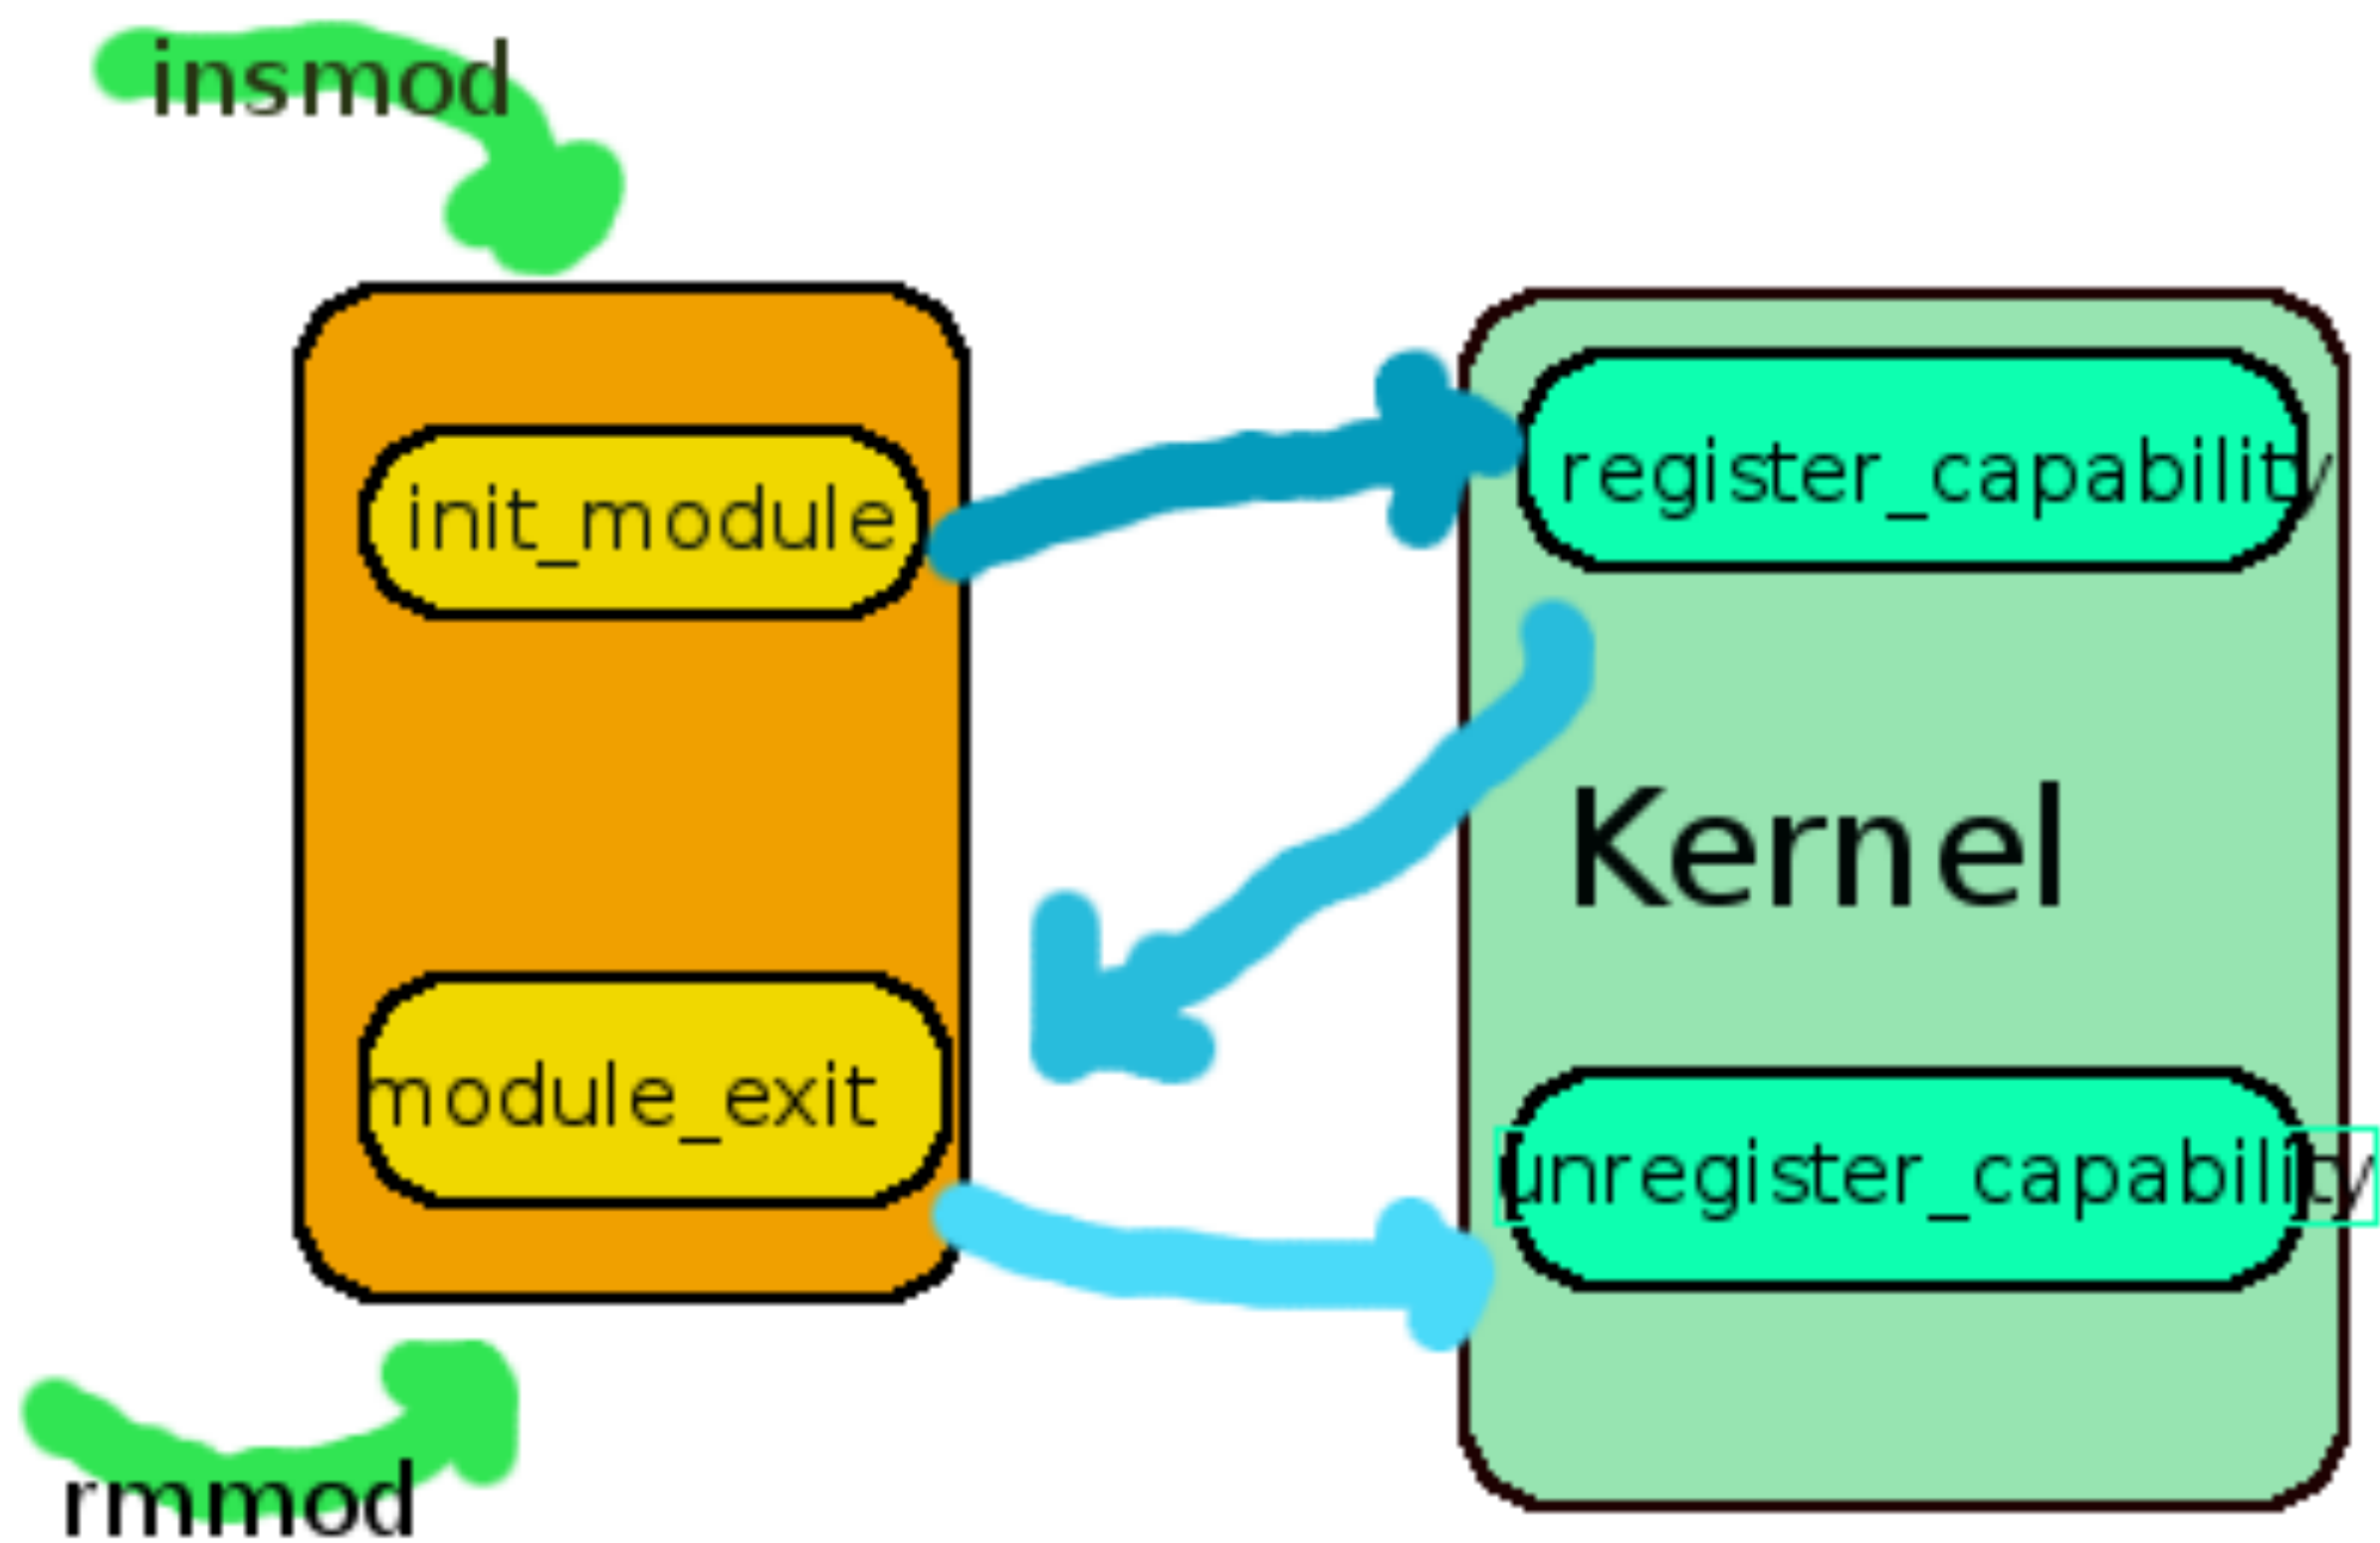
bridges

logical

hardware
interfaces

electronics





Caution!

- No standard C Lib
 - no printf, scanf, gets, etc..
 - no gdb (or any debugging tool)
- Use kernel functions only
 - printk

printk

- Usage

```
printk( [LEVEL] const char *format_string, ... );
```

- Example

```
printk(KERN_ERR "something went wrong, return  
code: %d\n", ret);
```

Level:

KERN_EMERG

KERN_ALERT

KERN_CRIT

KERN_ERR

KERN_WARNING

KERN_NOTICE

KERN_INFO

KERN_DEBUG

KERN_DEFAULT

KERN_CONT

Let's build
our first Kernel Module

Dummy Kernel Module

```
#include <linux/module.h>      /* Needed by all modules */
#include <linux/kernel.h>      /* Needed for KERN_INFO */

MODULE_LICENSE("GPL");
MODULE_AUTHOR("KRERK PIROMSOPA, PH.D. <Krerk.P@chula.ac.th>");
MODULE_DESCRIPTION("CP OS API");

int init_module(void)
{
    printk(KERN_INFO "CUCPMOD: init\n");
    /*
     * non 0 - means init_module failed
     *         - module can't be loaded.
     */
    return 0;
}

void cleanup_module(void)
{
    printk(KERN_INFO "CUCPMOD: cleanup\n");
}
```

CPMOD.C

Makefile

```
obj-m += cpmmod.o
```

```
all:
```

```
    make -C /lib/modules/$(shell uname -r)/build M=$(PWD) modules
```

```
clean:
```

```
    make -C /lib/modules/$(shell uname -r)/build M=$(PWD) clean
```


About Debian

- Don't forget to install
linux-kernel-headers
(linux-libc-dev)

Common Commands

- `lsmod` - Show the status of modules in the Linux Kernel
- `insmod` - Simple program to insert a module into the Linux Kernel
`insmod [filename] [module options...]`
- `rmmod` - Simple program to remove a module from the Linux Kernel
`rmmod [-f] [-s] [-v] [module name]`
- `modprobe` - Add and remove modules from the Linux Kernel
`modprobe [-v] [-V] [-C config-file] [-n] [-i] [-q] [-b] [module name]
[module parameters...]`
- `modinfo` - Show information about a Linux Kernel module
`modinfo [-o] [-F field] [-k kernel] [module name|filename...]`
- `dmesg` - print or control the kernel ring buffer



Device Drivers

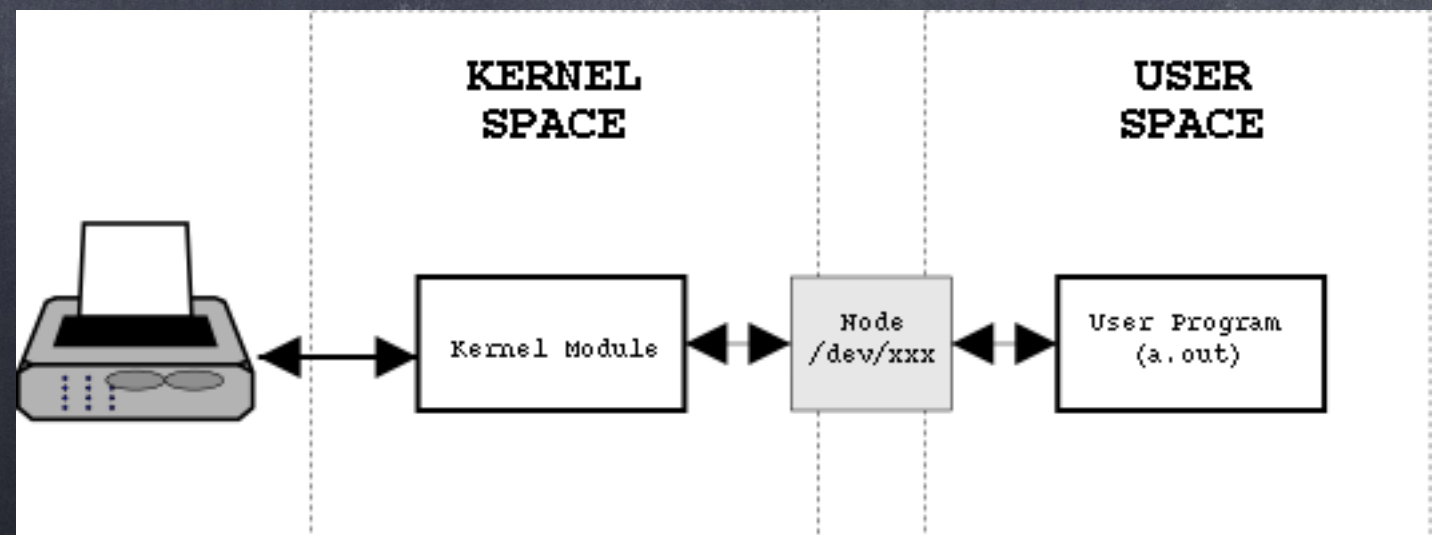


Linux Devices

- Device File
 - an interface for a device driver in a file system
(Like an ordinary file)
 - Use file permissions for device permissions
 - Use file operations for accessing device

Example

- A write to /dev/lpo will print to an associated printer



Just for fun..

- Open a terminal.
- Use "who" command to figure out the pseudo terminal device file.
- Try "echo data > /dev/pts/0"

```
krerk@OSBox:~$ who
krerk      pts/0      2016-11-22 15:23 (???.???.???.??)
krerk@OSBox:~$ echo "I am a master." >/dev/pts/0
I am a master.
krerk@OSBox:~$
```


Character Devices

- sending and receiving single characters
- Cannot seek
- eg. Serial/Parallel Port

Block Devices

- sending entire blocks of data
- Can seek
- eg. Disk, USB Camera

Major and Minor Device No.

- Devices are divided into sets called major device numbers.
- Same major number means same driver

```
brw-rw---- 1 root disk 8, 0 Nov 12 14:18 /dev/sda
brw-rw---- 1 root disk 8, 1 Nov 12 14:18 /dev/sda1
brw-rw---- 1 root disk 8, 2 Nov 12 14:18 /dev/sda2
brw-rw---- 1 root disk 8, 5 Nov 12 14:18 /dev/sda5
crw----- 1 kerk tty 136, 0 Nov 23 22:32 /dev/pts/0
```


To create a device

- Try
mknod /dev/osinfo c 250 0
- This will create

```
crw-r--r-- 1 root root 250, 0 Nov 23 06:51 /dev/osinfo
```


File Operations

```
struct file_operations {
    struct module *owner;
    loff_t(*llseek) (struct file *, loff_t, int);
    ssize_t(*read) (struct file *, char __user *, size_t, loff_t *);
    ssize_t(*aio_read) (struct kiocb *, char __user *, size_t, loff_t);
    ssize_t(*write) (struct file *, const char __user *, size_t, loff_t *);
    ssize_t(*aio_write) (struct kiocb *, const char __user *, size_t, loff_t);
    int (*readdir) (struct file *, void *, filldir_t);
    unsigned int (*poll) (struct file *, struct poll_table_struct *);
    int (*ioctl) (struct inode *, struct file *, unsigned int, unsigned long);
    int (*mmap) (struct file *, struct vm_area_struct *);
    int (*open) (struct inode *, struct file *);
    int (*flush) (struct file *);
    int (*release) (struct inode *, struct file *);
    int (*fsync) (struct file *, struct dentry *, int datasync);
    int (*aio_fsync) (struct kiocb *, int datasync);
    int (*fasync) (int, struct file *, int);
    int (*lock) (struct file *, int, struct file_lock *);
    ssize_t(*readv) (struct file *, const struct iovec *, unsigned long, loff_t *);
    ssize_t(*writev) (struct file *, const struct iovec *, unsigned long, loff_t *);
    ssize_t(*sendfile) (struct file *, loff_t *, size_t, read_actor_t, void __user *);
    ssize_t(*sendpage) (struct file *, struct page *, int, size_t, loff_t *, int);
    unsigned long (*get_unmapped_area) (struct file *, unsigned long, unsigned long, unsigned long,
    unsigned long);
};
```

see `linux/fs.h` for more details

TODO

- Implement only related operations
- Leave others to NULL


A better way

- Use C99 syntax for convenient

```
struct file_operations fops = {  
    .read = device_read,  
    .write = device_write,  
    .open = device_open,  
    .release = device_release  
};
```


To register device

- `int register_chrdev(unsigned int major, const char *name, struct file_operations *fops);`
- `unregister_chrdev(unsigned int major, const char *name);`



see `linux/fs.h` for more details

Example: Read-only char device

```
static int device_open(struct inode *, struct file *);  
static int device_release(struct inode *inode, struct file *file);  
static ssize_t device_read(struct file *, char *, size_t, loff_t *);
```

```
// File operations structor  
// Only implement those that will be used.  
static struct file_operations dev_fops = {  
    .read = device_read,  
    .open = device_open,  
    .release = device_release  
};
```


More about C

- For security, try to use static modifiers for functions, variables in C to make it local to your kernel module.

More about driver

- To avoid module from being unloaded (rmmod) while using,
 - On every access to module (e.g. device_open), please lock the module with
`try_module_get(THIS_MODULE);`
 - On release, release the module with
`module_put(THIS_MODULE);`

Let's have fun.

Q1. Dummy Mod

- Make a dummy driver that will do the following:
 - On initial, display
KERN_INFO "CPMOD: init"
 - On cleanup, display
KERN_INFO "CPMOD: cleanup\n"

Q2. Char Driver

- Create osinfo driver (minor no. 0) (at /dev/osinfoa).
- Reading from this device file will display:

```
0:CP ENG CU OS 2018S1 - Instructors
1:      Krerk Piromsopa, Ph. D.
2:      Kunwadee Sripanidkulchai, Ph. D.
3:      Thongchai Rojkangsadan
4:      Veera Muangsin, Ph.D.
```


Q3. Adv. Char Driver

- Modify driver in Q2 to display to detect minor no.
- If minor no. == 1, reading from the device will display names of your team members.

0:CP ENG CU OS 2018S1 – Students
1: Kerk Piromsopa, Ph. D.

Hint: To detect minor number, use
`static inline unsigned iminor(const struct inode *inode)`

Q4. CPU model driver

- Use the following function (taken from arch/x86/include/asm/processor.h) to create a device driver for extracting cpuid (vendor ID, features, serial no.).

```
static inline void native_cpuid(unsigned int *eax, unsigned int *ebx,
                                unsigned int *ecx, unsigned int *edx)
{
    /* ecx is often an input as well as an output. */
    asm volatile("cpuid"
        : "=a" (*eax),
          "=b" (*ebx),
          "=c" (*ecx),
          "=d" (*edx)
        : "0" (*eax), "2" (*ecx)
        : "memory");
}
```

```
eax = 3; /* processor serial number */
native_cpuid(&eax, &ebx, &ecx, &edx);
printf("serial number 0x%08x%08x\n", edx, ecx);
```

For more details , see <https://en.wikipedia.org/wiki/CPUID>

More Hint.

- Kernel supports limited string function.
- For Q4, try
 - `snprintf`, `strncat`.
- See Kernel API (<https://www.kernel.org/doc/htmldocs/kernel-api/>) for more details.

Quick demonstration

References

- <https://www.kernel.org/>
- Kernel API <https://www.kernel.org/doc/htmldocs/kernel-api/>
- <http://elixir.free-electrons.com/linux/latest/source>