


Chapter 5: Process Scheduling





Chapter 5: Process Scheduling

- Basic Concepts
- Scheduling Criteria
- Scheduling Algorithms
- Thread Scheduling
- Multiple-Processor Scheduling
- Real-Time CPU Scheduling
- Operating Systems Examples
- Algorithm Evaluation





Objectives

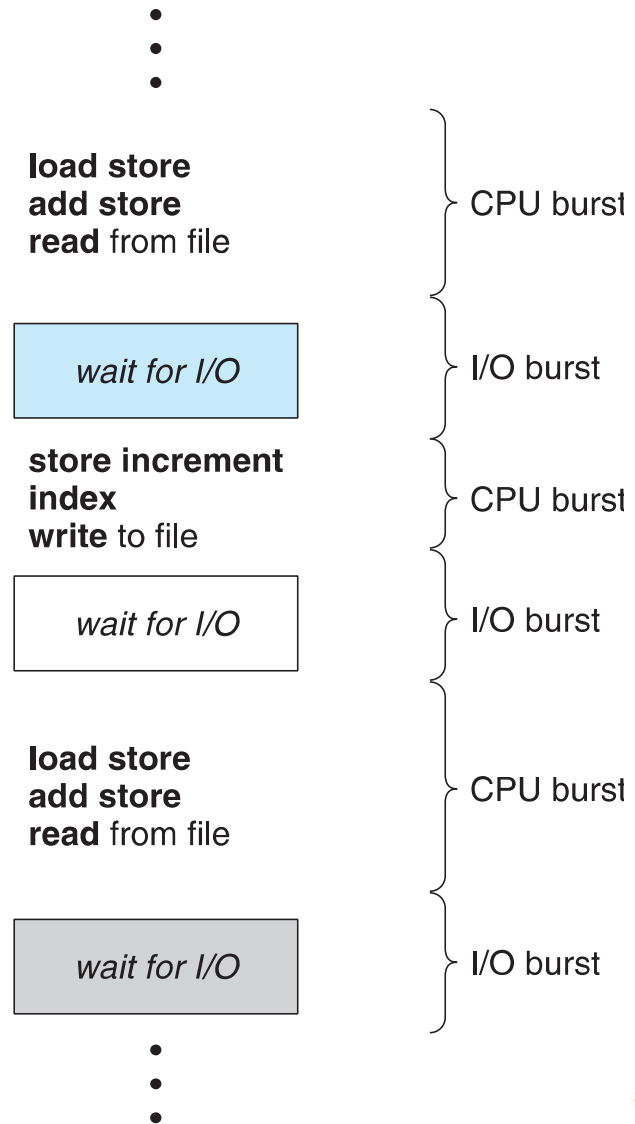
- To introduce CPU scheduling, which is the basis for multiprogrammed operating systems
- To describe various CPU-scheduling algorithms
- To discuss evaluation criteria for selecting a CPU-scheduling algorithm for a particular system
- To examine the scheduling algorithms of several operating systems






Basic Concepts

- ❑ Maximum CPU utilization obtained with multiprogramming
- ❑ CPU-I/O Burst Cycle – Process execution consists of a **cycle** of CPU execution and I/O wait
- ❑ **CPU burst** followed by **I/O burst**
- ❑ CPU burst distribution is of main concern





Histogram of CPU-burst Times





CPU Scheduler

- **Short-term scheduler** selects from among the processes in ready queue, and allocates the CPU to one of them
 - Queue may be ordered in various ways
- CPU scheduling decisions may take place when a process:
 1. Switches from running to waiting state
 2. Switches from running to ready state
 3. Switches from waiting to ready
 4. Terminates
- Scheduling under 1 and 4 is **nonpreemptive**
- All other scheduling is **preemptive**
 - Consider access to shared data
 - Consider preemption while in kernel mode
 - Consider interrupts occurring during crucial OS activities





Dispatcher

- Dispatcher module gives control of the CPU to the process selected by the short-term scheduler; this involves:
 - switching context
 - switching to user mode
 - jumping to the proper location in the user program to restart that program

- **Dispatch latency** – time it takes for the dispatcher to stop one process and start another running





Scheduling Criteria

- **CPU utilization** – keep the CPU as busy as possible
- **Throughput** – # of processes that complete their execution per time unit
- **Turnaround time** – amount of time to execute a particular process
- **Waiting time** – amount of time a process has been waiting in the ready queue
- **Response time** – amount of time it takes from when a request was submitted until the first response is produced, not output (for time-sharing environment)





Scheduling Algorithm Optimization Criteria

- Max CPU utilization
- Max throughput
- Min turnaround time
- Min waiting time
- Min response time





First-Come, First-Served (FCFS) Scheduling

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3

- Suppose that the processes arrive in the order: P_1, P_2, P_3
The Gantt Chart for the schedule is:



- Waiting time for $P_1 = 0$; $P_2 = 24$; $P_3 = 27$
- Average waiting time: $(0 + 24 + 27)/3 = 17$
- Convoy effect** - short process behind long process
 - Consider one CPU-bound and many I/O-bound processes





FCFS Scheduling (Cont.)

Suppose that the processes arrive in the order:

P_2, P_3, P_1

□ The Gantt chart for the schedule is:



- Waiting time for $P_1 = 6$; $P_2 = 0$; $P_3 = 3$
- Average waiting time: $(6 + 0 + 3)/3 = 3$
- Much better than previous case





Shortest-Job-First (SJF) Scheduling

- Associate with each process the length of its next CPU burst
 - Use these lengths to schedule the process with the shortest time

- SJF is **optimal** – gives minimum average waiting time for a given set of processes
 - The difficulty is knowing the length of the next CPU request
 - Could predict the next CPU burst based on previous ones.






Example of SJF

<u>Process</u>	<u>Burst Time</u>
P_1	6
P_2	8
P_3	7
P_4	3

- SJF scheduling chart



- Average waiting time = $(3 + 16 + 9 + 0) / 4 = 7$






Example of Shortest-remaining-time-first

- Now we add the concepts of varying arrival times and preemption to the analysis

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

- Preemptive SJF Gantt Chart*



- Average waiting time = $[(10-1)+(1-1)+(17-2)+(5-3)]/4 = 26/4 = 6.5$ msec





Priority Scheduling

- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority (smallest integer \equiv highest priority)
 - Preemptive
 - Nonpreemptive
- SJF is priority scheduling where priority is the inverse of predicted next CPU burst time
- Problem \equiv **Starvation** – low priority processes may never execute
- Solution \equiv **Aging** – as time progresses increase the priority of the process





Example of Priority Scheduling

<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>
P_1	10	3
P_2	1	1
P_3	2	4
P_4	1	5
P_5	5	2

□ Priority scheduling Gantt Chart



□ Average waiting time = 8.2 msec





Round Robin (RR)

- Each process gets a small unit of CPU time (**time quantum** q), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
- If there are n processes in the ready queue and the time quantum is q , then each process gets $1/n$ of the CPU time in chunks of at most q time units at once. No process waits more than $(n-1)q$ time units.
- Timer interrupts every quantum to schedule next process
- Performance
 - q large \Rightarrow FIFO
 - q small \Rightarrow q must be large with respect to context switch, otherwise overhead is too high






Example of RR with Time Quantum = 4

Process	Burst Time
P_1	24
P_2	3
P_3	3

□ The Gantt chart is:




- Typically, higher average turnaround than SJF, but better **response**
- q should be large compared to context switch time
- q usually 10ms to 100ms, context switch < 10 usec






Time Quantum and Context Switch Time





Turnaround Time Varies With The Time Quantum



process	time
P_1	6
P_2	3
P_3	1
P_4	7

80% of CPU bursts should be shorter than q





Multilevel Queue

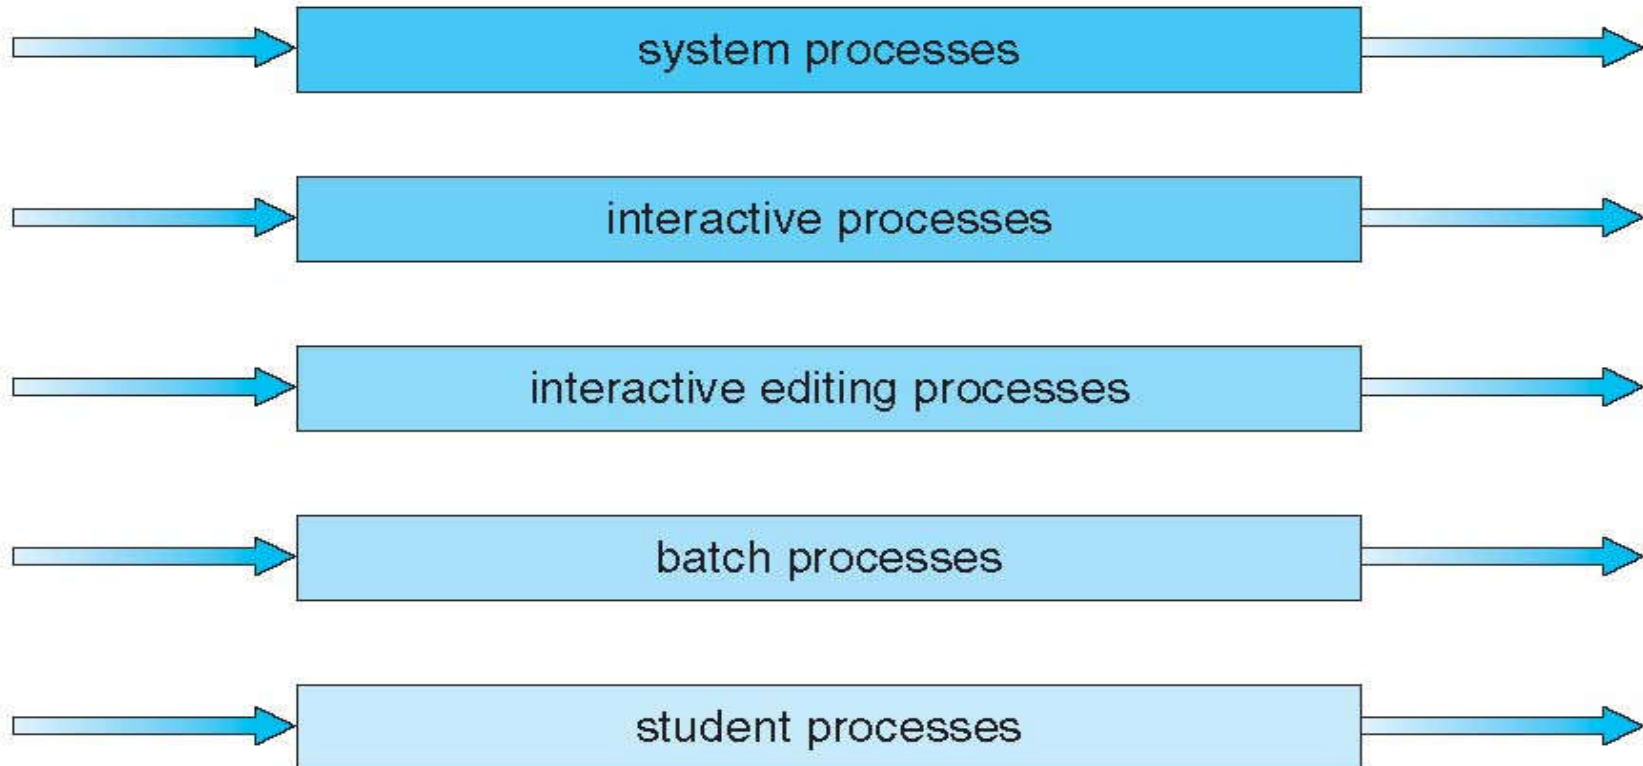
- Ready queue is partitioned into separate queues, eg:
 - **foreground** (interactive)
 - **background** (batch)
- Process permanently in a given queue
- Each queue has its own scheduling algorithm:
 - foreground – RR
 - background – FCFS
- Scheduling must be done between the queues:
 - Fixed priority scheduling; (i.e., serve all from foreground then from background). Possibility of starvation.
 - Time slice – each queue gets a certain amount of CPU time which it can schedule amongst its processes; i.e., 80% to foreground in RR
 - 20% to background in FCFS





Multilevel Queue Scheduling

highest priority



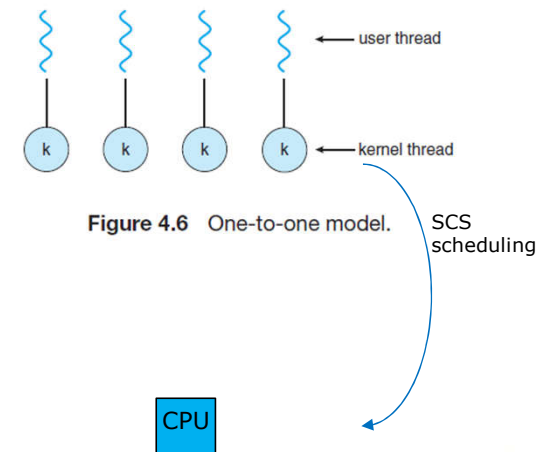
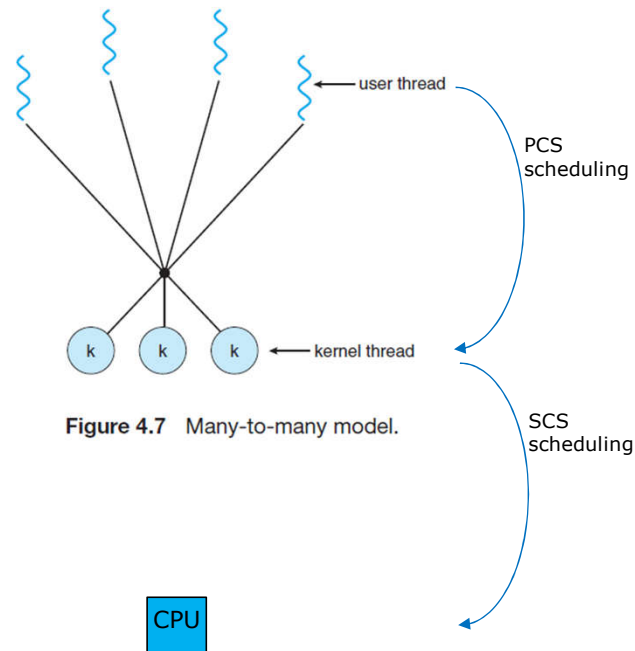
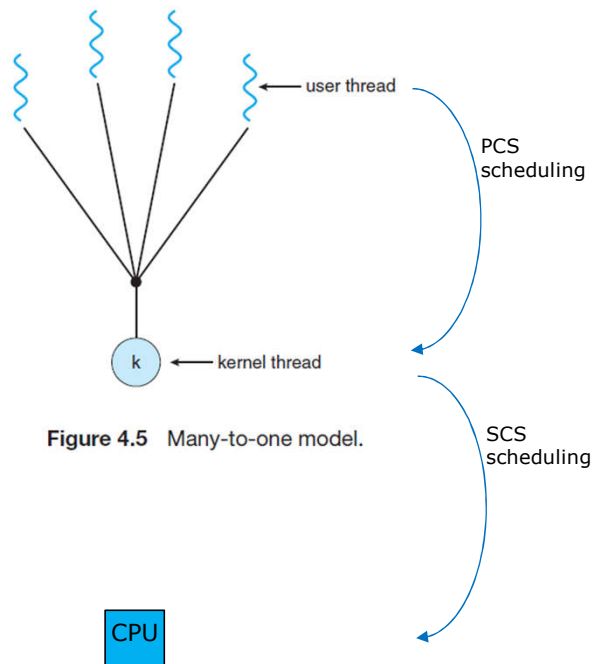
lowest priority





Thread Scheduling


- **Process-contention scope (PCS)** – scheduling among threads belonging to the same process
- **System-contention scope (SCS)** – scheduling among all threads in the system
- Systems using one-to-one model (e.g. Windows, Linux) uses only SCS.





Multiple-Processor Scheduling

- With multi-core/multi-processor, parallel computing and load sharing become possible.
- Each processor has its own private queue of ready processes.
- Processor Affinity
 - Keep each process/thread running on the same processor.
 - To benefit from cache memory.
- Load balancing
 - Keep workload evenly distributed across all processors.
 - To benefit from multiple-processor.
 - Migrate processes from overloaded processors to less-busy processors.
- Load balancing often counteracts the benefits of process affinity.





DEMONSTRATION





To see every process on the system using standard syntax:

```
ps -e
ps -ef
ps -eF
ps -ely
```

To see every process on the system using BSD syntax:

```
ps ax
ps axu
```

To print a process tree:

```
ps -ejH
ps axjf
```

To get info about threads:

```
ps -eLf
ps axms
```

```
pom@X280:/mnt/c/Users/Pom$ ps axu
```

USER	PID	%CPU	%MEM	VSZ	RSS	TTY	STAT	START	TIME	COMMAND
root	1	0.0	0.0	1744	432	?	Sl	Feb17	0:00	/init
root	111	0.0	0.0	1764	0	?	Ss	Feb17	0:00	/init
root	112	0.0	0.0	1764	116	?	S	Feb17	0:00	/init
root	113	0.0	0.5	1088556	10428	pts/0	Ssl+	Feb17	0:21	/mnt/wsl/docker-desktop/docker-desktop-proxy --distro-name Ubuntu-20.04
root	119	0.0	0.0	1764	4	?	S	Feb17	0:00	/init
pom	120	0.0	0.7	763732	14736	pts/1	Ssl+	Feb17	0:36	docker serve --address unix:///home/pom/.docker/run/docker-cli-api.sock
root	146	0.0	0.0	1764	40	?	Ss	17:24	0:00	/init
root	147	0.0	0.0	1764	64	?	S	17:24	0:00	/init
pom	148	0.0	0.2	10188	5300	pts/2	Ss	17:24	0:00	-bash
pom	389	0.0	0.1	10604	3300	pts/2	R+	18:11	0:00	ps axu

```
pom@X280:/mnt/c/Users/Pom$ ps axjf
```

PPID	PID	PGID	SID	TTY	TPGID	STAT	UID	TIME	COMMAND
0	1	0	0	?	-1	Sl	0	0:00	/init
1	111	111	111	?	-1	Ss	0	0:00	/init
111	112	111	111	?	-1	S	0	0:00	_ /init
112	113	113	113	pts/0	113	Ssl+	0	0:21	_ /mnt/wsl/docker-desktop/docker-desktop-proxy --distro-name Ubuntu-20.04
111	119	111	111	?	-1	S	0	0:00	_ /init
119	120	120	120	pts/1	120	Ssl+	1000	0:36	_ docker serve --address unix:///home/pom/.docker/run/docker-cli-api.sock
1	146	146	146	?	-1	Ss	0	0:00	/init
146	147	146	146	?	-1	S	0	0:00	_ /init
147	148	148	148	pts/2	397	Ss	1000	0:00	_ -bash
148	397	397	148	pts/2	397	R+	1000	0:00	_ ps axjf





- Information about process is kept in /proc/<pid>

```
pom@X280:/mnt/c/Users/Pom$ ls /proc
1    146    bus    crypto    filesystems    kallsyms    kpagecount    misc    partitions    swaps    uptime
111  147    cgroups    devices    fs    kcore    kpageflags    modules    sched_debug    sys    version
112  148    cmdline    diskstats    interrupts    key-users    loadavg    mounts    schedstat    sysvipc    vmallocinfo
113  401    config.gz    dma    iomem    keys    locks    mtrr    self    thread-self    vmstat
119  acpi    consoles    driver    ioports    kmsg    mdstat    net    softirqs    timer_list    zoneinfo
120  buddyinfo    cpuinfo    execdomains    irq    kpagecgroup    meminfo    pagetypeinfo    stat    tty
```

```
pom@X280:/mnt/c/Users/Pom$ ls /proc/148
arch_status    cmdline    environ    io    mountinfo    oom_adj    projid_map    smaps    status    uid_map
attr           comm       exe        limits    mounts    oom_score    root    smaps_rollup    syscall    wchan
auxv           coredump_filter    fd        map_files    mountstats    oom_score_adj    sched    stack    task
cgroup         cpuset     fdinfo     maps      net        pagemap    schedstat    stat    timers
clear_refs     cwd        gid_map    mem       ns         personality    setgroups    statm    timerslack_ns
```





Context Switches

Run “top” command in one console.

```
pom@X280:/mnt/c/Users/Pom$ top -d 1
```

Update every 1 second

```
top - 18:50:33 up 4 days, 12:32, 0 users, load average: 0.24, 0.28, 0.31
Tasks: 13 total, 1 running, 12 sleeping, 0 stopped, 0 zombie
%Cpu(s): 1.0 us, 2.0 sy, 0.0 ni, 96.5 id, 0.0 wa, 0.0 hi, 0.5 si, 0.0 st
MiB Mem : 1916.6 total, 104.7 free, 1230.5 used, 581.4 buff/cache
MiB Swap: 1024.0 total, 229.8 free, 794.2 used. 366.1 avail Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
1	root	20	0	1744	432	396	S	0.0	0.0	0:00.02	init
111	root	20	0	1764	0	0	S	0.0	0.0	0:00.00	init
112	root	20	0	1764	116	116	S	0.0	0.0	0:00.02	init
113	root	20	0	1088556	11988	3644	S	0.0	0.6	0:21.43	docker-desktop-
119	root	20	0	1764	4	0	S	0.0	0.0	0:00.00	init
120	pom	20	0	763732	14736	0	S	0.0	0.8	0:36.52	docker
146	root	20	0	1764	40	0	S	0.0	0.0	0:00.00	init
147	root	20	0	1764	64	0	S	0.0	0.0	0:00.39	init
148	pom	20	0	10188	5300	3464	S	0.0	0.3	0:00.57	bash
437	root	20	0	1764	40	0	S	0.0	0.0	0:00.00	init
438	root	20	0	1764	64	0	S	0.0	0.0	0:00.02	init
439	pom	20	0	10056	5148	3408	S	0.0	0.3	0:00.11	bash
485	pom	20	0	10860	3720	3216	R	0.0	0.2	0:00.10	top

Run these commands in another console.

```
pom@X280:/mnt/c/Users/Pom$ date; grep ctxt /proc/485/status
Mon Feb 21 18:49:48 +07 2022
voluntary_ctxt_switches:      90
nonvoluntary_ctxt_switches:    1
pom@X280:/mnt/c/Users/Pom$ date; grep ctxt /proc/485/status
Mon Feb 21 18:49:54 +07 2022
voluntary_ctxt_switches:      97
nonvoluntary_ctxt_switches:    1
pom@X280:/mnt/c/Users/Pom$ date; grep ctxt /proc/485/status
Mon Feb 21 18:50:02 +07 2022
voluntary_ctxt_switches:      104
nonvoluntary_ctxt_switches:    1
```

This output shows the number of context switches over the lifetime of the process. Notice the distinction between *voluntary* and *nonvoluntary* context switches. A voluntary context switch occurs when a process has given up control of the CPU because it requires a resource that is currently unavailable (such as blocking for I/O.) A nonvoluntary context switch occurs when the CPU has been taken away from a process, such as when its time slice has expired or it has been preempted by a higher-priority process.





```
pom@X280:/mnt/c/Users/Pom$ top -d 0.01
```

Update every 10 msec

```
pom@X280:/mnt/c/Users/Pom$ date; grep ctxt /proc/656/status
Mon Feb 21 19:40:54 +07 2022
voluntary_ctxt_switches:      3784
nonvoluntary_ctxt_switches:   742
pom@X280:/mnt/c/Users/Pom$ date; grep ctxt /proc/656/status
Mon Feb 21 19:41:03 +07 2022
voluntary_ctxt_switches:      4681
nonvoluntary_ctxt_switches:   849
pom@X280:/mnt/c/Users/Pom$
```

Nonvoluntary context switch about every 100 msec.

```
pom@X280:~/OS$ cat infinite.c
void main()
{
    for (;;)
        ;
}
```

```
pom@X280:~/OS$ ./infinite
```

```
pom@X280:/mnt/c/Users/Pom$ date; grep ctxt /proc/730/status
Mon Feb 21 20:34:18 +07 2022
voluntary_ctxt_switches:      0
nonvoluntary_ctxt_switches:   1585
pom@X280:/mnt/c/Users/Pom$ date; grep ctxt /proc/730/status
Mon Feb 21 20:34:28 +07 2022
voluntary_ctxt_switches:      0
nonvoluntary_ctxt_switches:   1926
```

Nonvoluntary context switch about every 30 msec.





Nice (set priority of process)

```
NICE(1)                                User Commands                                NICE(1)

NAME
    nice - run a program with modified scheduling priority

SYNOPSIS
    nice [OPTION] [COMMAND [ARG]...]

DESCRIPTION
    Run COMMAND with an adjusted niceness, which affects process scheduling. With no COMMAND, print the current niceness. Niceness values range from -20 (most favorable to the process) to 19 (least favorable to the process).

    Mandatory arguments to long options are mandatory for short options too.

    -n, --adjustment=N
        add integer N to the niceness (default 10)
```

```
pom@X280:~/OS$ cat multiply.c
int main(void) {
    int total = 1;
    for (int j =1; j <=50000 ; j++)
        for (int k =1; k <=50000 ; k++)
            total *= j*k;
}
```

```
pom@X280:~/OS$ cat multiply.sh
time ./multiply &
time nice -n 1 ./multiply &
time nice -n 2 ./multiply &
time nice -n 3 ./multiply &
```

```
pom@X280:~/OS$ ./multiply.sh
pom@X280:~/OS$
real    0m14.801s
user    0m8.918s
sys     0m0.041s

real    0m19.142s
user    0m9.243s
sys     0m0.067s

real    0m21.373s
user    0m9.237s
sys     0m0.059s

real    0m22.535s
user    0m8.989s
sys     0m0.086s
```

```
top - 23:29:09 up 4 days, 17:11,  0 users,  load average: 0.92, 0.53, 0.47
Tasks:  21 total,   5 running, 16 sleeping,   0 stopped,   0 zombie
%Cpu(s):  97.8/2.2  100[|||||||||||||||||||||||||||||||||||||||||||||||||
MiB Mem :  1916.6 total,    73.2 free,  1250.4 used,   593.0 buff/cache
MiB Swap:  1024.0 total,   290.5 free,   733.5 used.  407.4 avail Mem

  PID USER      PR  NI   VIRT   RES   SHR  S  %CPU  %MEM    TIME+  COMMAND
  860 pom       20   0   2360    580   516  R   62.5   0.0   0:05.47  multiply
  864 pom       21   1   2360    576   512  R   49.2   0.0   0:04.28  multiply
  863 pom       22   2   2360    584   520  R   38.9   0.0   0:03.52  multiply
  865 pom       23   3   2360    516   452  R   30.6   0.0   0:02.76  multiply
```

