


Chapter 9: Virtual-Memory Management






Virtual Memory

- Goal
 - Provide memory space for all processes that can be much larger than physical memory.
- How
 - Only some part of program code and data need to be in physical memory at a time.
 - ▶ principle of locality
 - The rest of program code and data can be stored in secondary storage (hard disk)
 - ▶ “swap partition” or “swap file”
 - ▶ extend memory hierarchy
 - Implement “Demand Paging” mechanism in OS
 - ▶ similar to “cache miss”
- Other benefits
 - Address spaces can be shared by several processes
 - ▶ shared program code, shared libraries, shared memory (IPC)
 - fast process creation
 - memory-mapped file, memory-mapped I/O





Virtual Memory





Demand Paging

- Bring a page into memory only when it is needed
 - Less I/O needed
 - Less memory needed
 - Faster response
 - More processes
- CPU instruction references to anywhere in a page \Rightarrow the page needs to be in physical memory
 - invalid reference \Rightarrow abort
 - not-in-memory \Rightarrow bring to memory
 - no-free-frame \Rightarrow swap





Page table that supports demand paging

0	A
1	B
2	C
3	D
4	E
5	F
6	G
7	H


logical
memory

valid-invalid bit		
frame		
0	4	v
1		i
2	6	v
3		i
4		i
5	9	v
6		i
7		i

page table

0	
1	
2	
3	
4	A
5	
6	C
7	
8	
9	F
10	
11	
12	
13	
14	
15	

physical memory





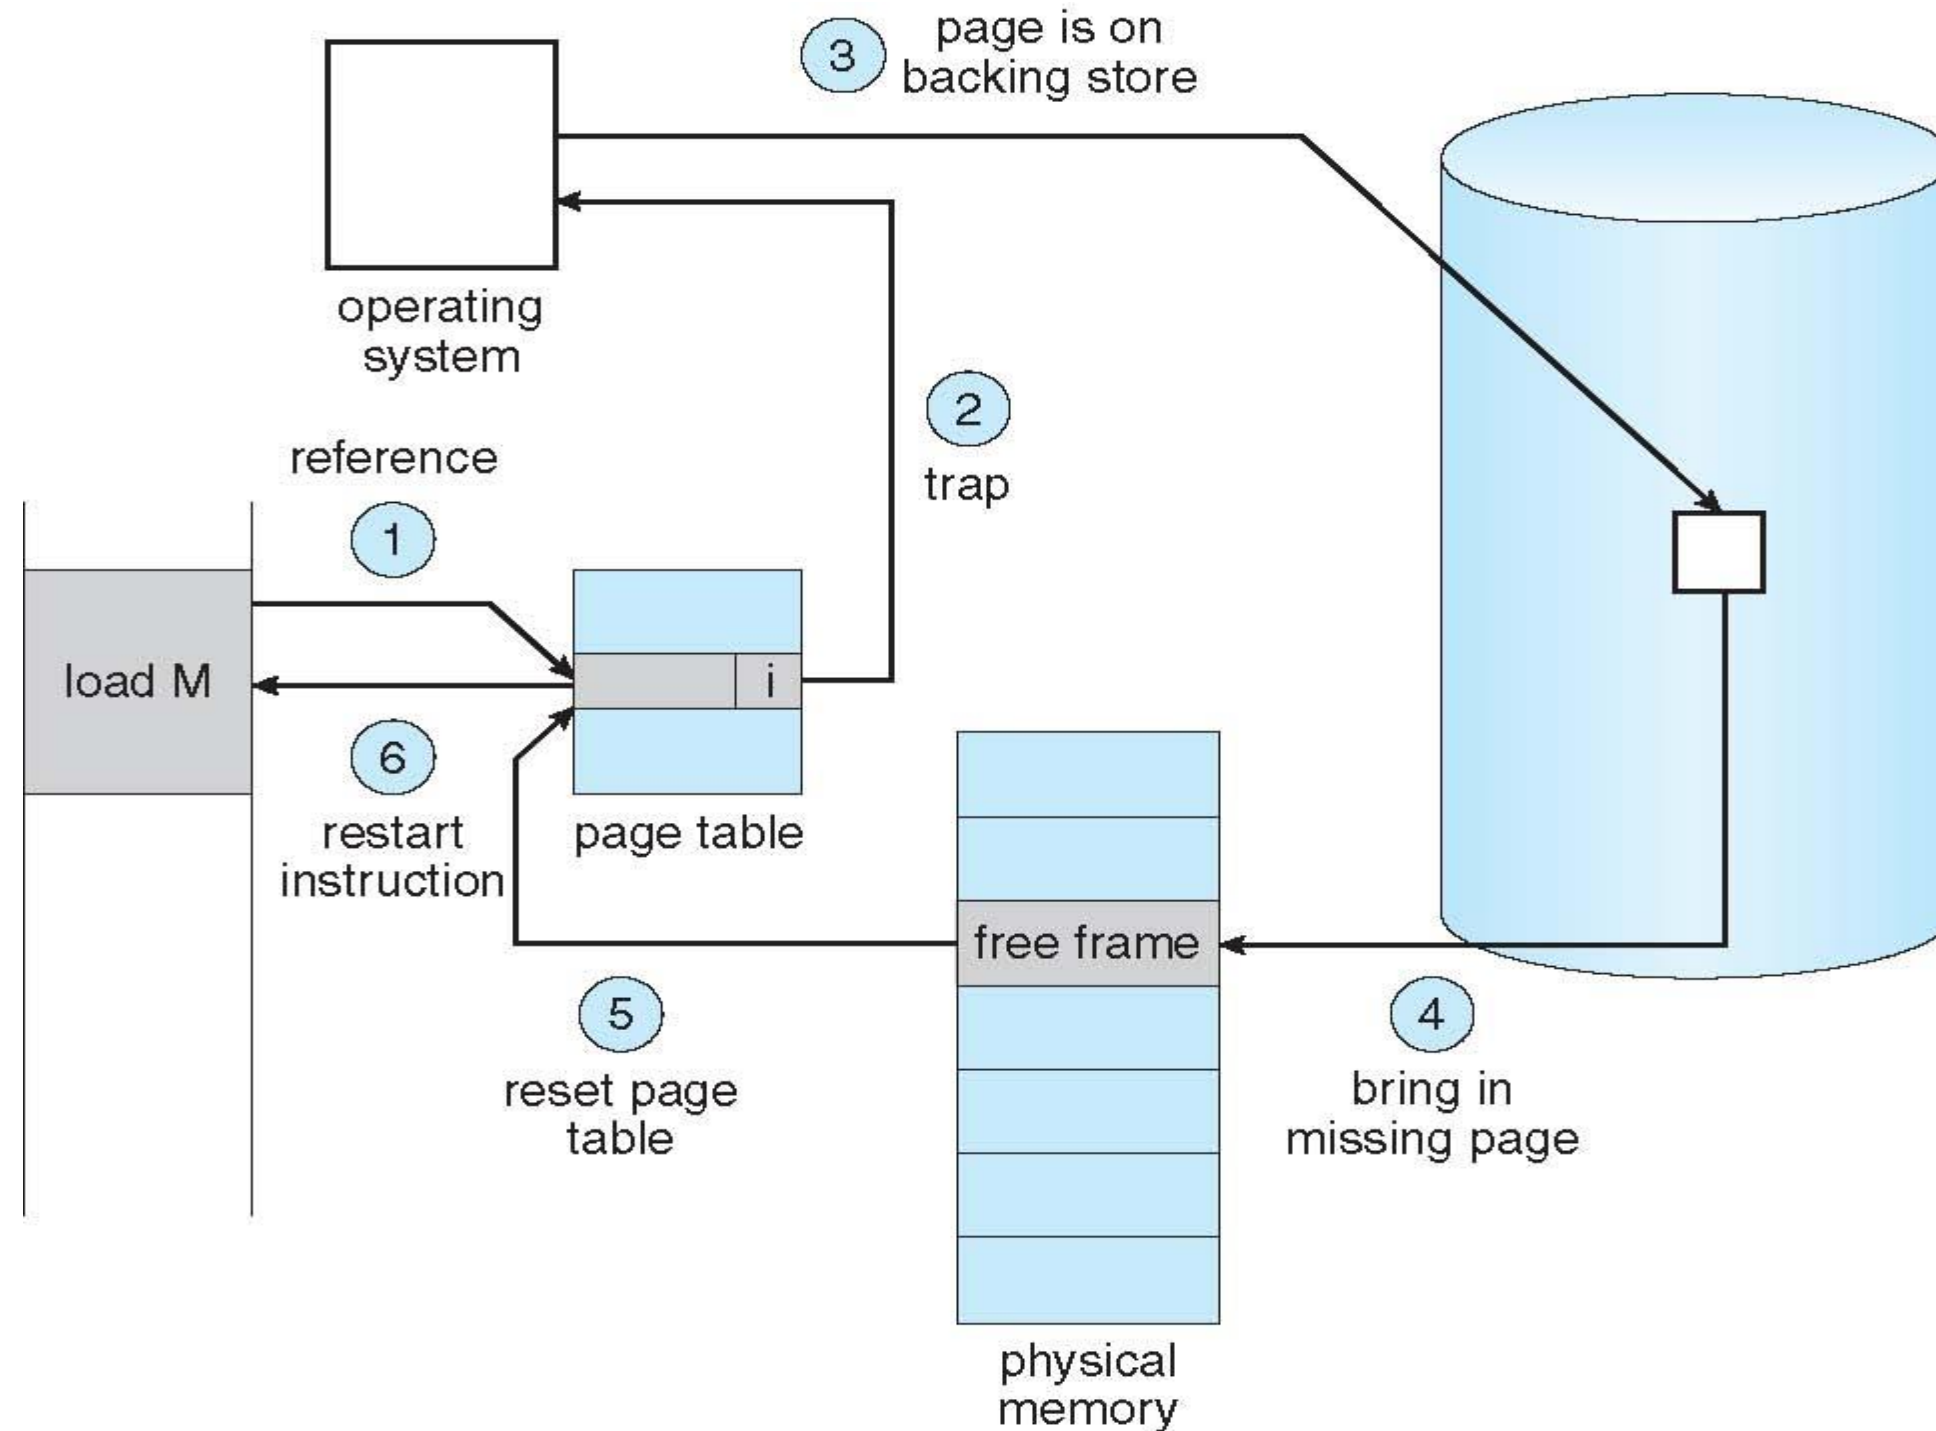
Page Fault

- A reference to an invalid page (e.g. first reference to that page) will cause an interrupt to operating system:
 - page fault interrupt
 - page fault handler
- 1. Operating system looks at another table to decide:
 - Invalid reference → abort
 - Just not in memory
- 2. Get empty frame
- 3. Swap page into frame
- 4. Update tables
- 5. Set validation bit = **v**
- 6. Restart the instruction that caused the page fault






Steps in Handling a Page Fault



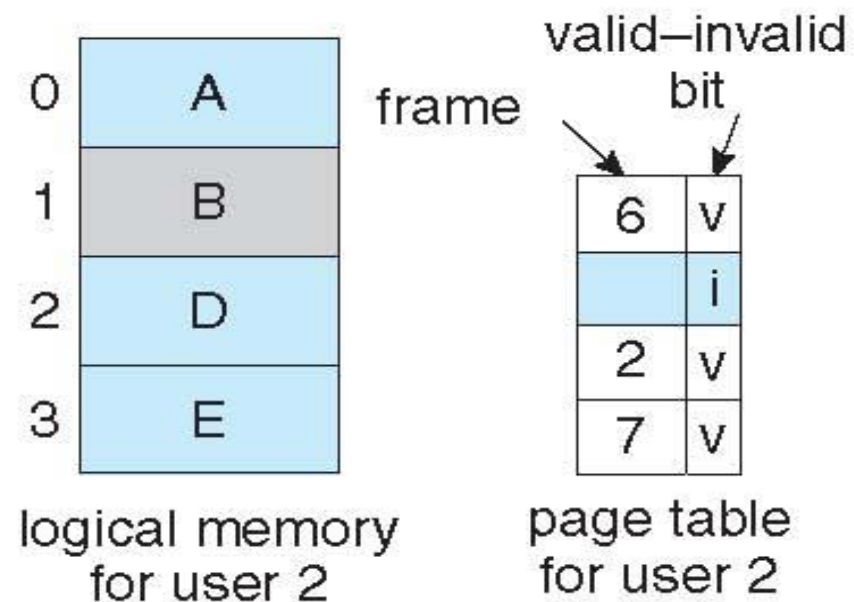




Need For Page Replacement



Program references to M in page 3 which is not in memory.
So, page fault occurs, without a free frame in physical memory.





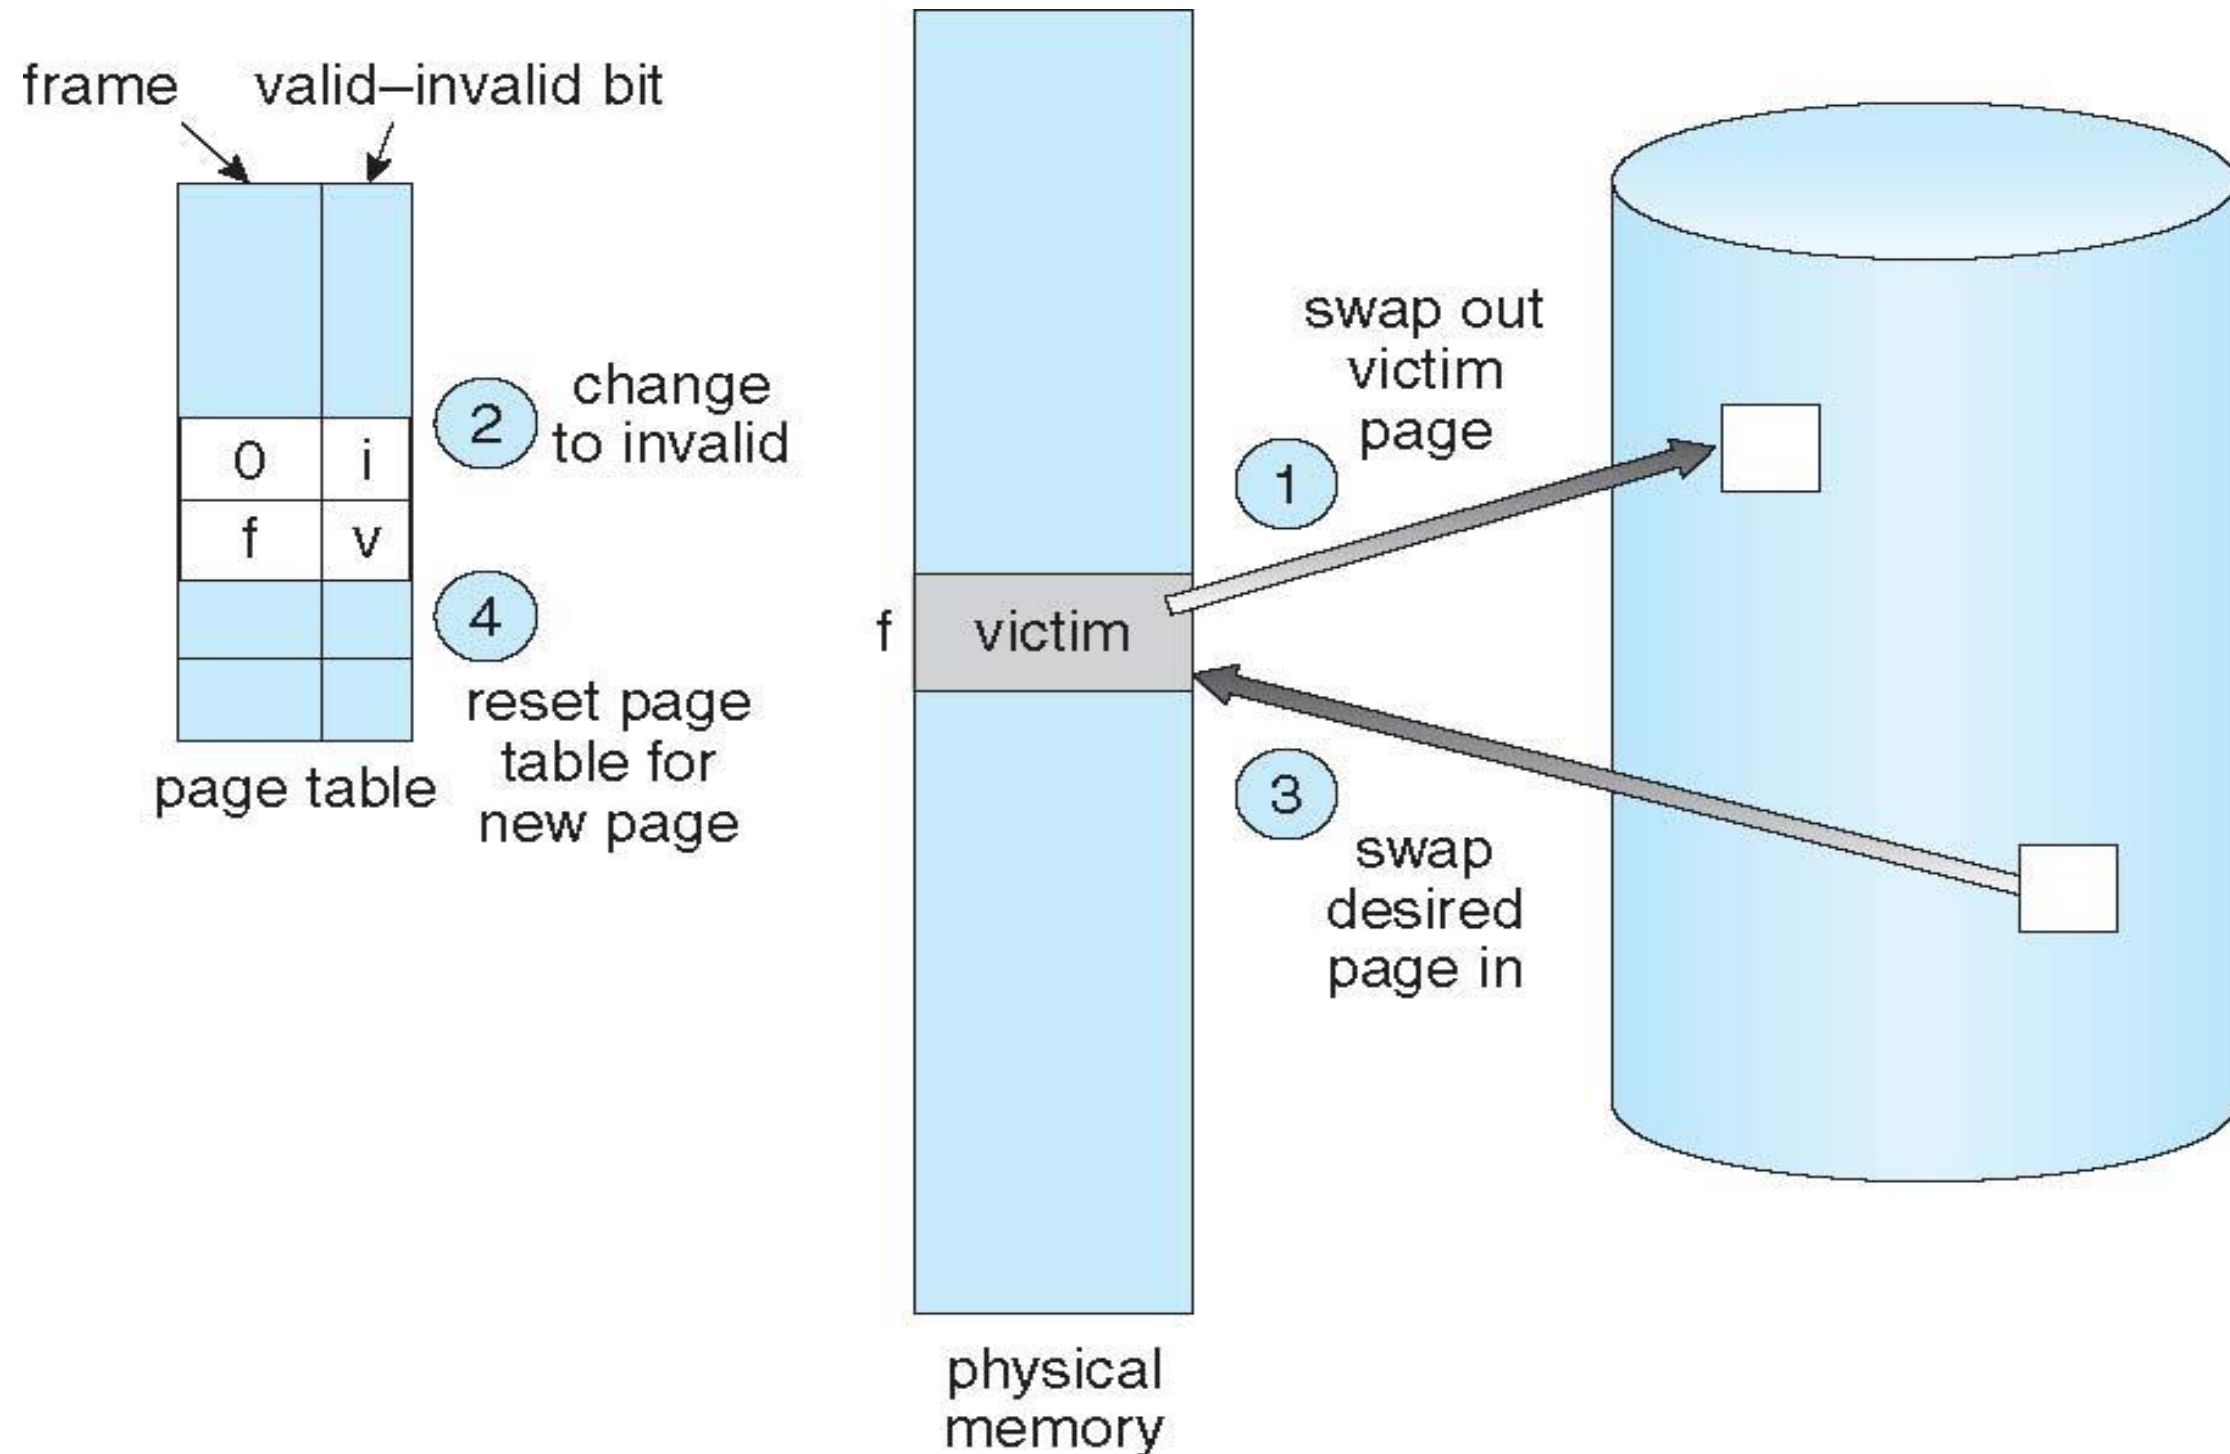
Basic Page Replacement

1. Find the location of the desired page on disk
2. Find a free frame:
 - If there is a free frame, use it
 - If there is no free frame, use a page replacement algorithm to select a **victim frame**
 - Write victim frame to disk if dirty
3. Bring the desired page into the (newly) free frame; update the page and frame tables
4. Continue the process by restarting the instruction that caused the trap






Page Replacement






Shared Library Using Virtual Memory





Memory-Mapped Shared Memory





Memory-Mapped Files

- Memory-mapped file I/O allows file I/O to be treated as routine memory access by **mapping** a disk block to a page in memory
- A file is initially read using demand paging
 - A page-sized portion of the file is read from the file system into a physical page
 - Subsequent reads/writes to/from the file are treated as ordinary memory accesses
- Simplifies and speeds file access by driving file I/O through memory rather than `read()` and `write()` system calls
- Also allows several processes to map the same file allowing the pages in memory to be shared
- But when does written data make it to disk?
 - Periodically and / or at file `close()` time





Memory Mapped Files

